

NFL

October 5, 2020

1 Goal

Build one or more regression models to determine the scores for each team using the other columns as features

Predict the scores, predict the outcome of the game

Do feature engineering. Build a few regression models.

Citing Sources Along with the sources I sprinkled in, I used these sources below throughout the text to better attack this problem.

<https://towardsdatascience.com/how-to-improve-sports-betting-odds-step-by-step-guide-in-python-94626b852f45>

```
[1]: import numpy as np
import pandas as pd
```

```
[2]: data = pd.read_csv('~/.Documents/EECS/EECS_731/HW/EECS731_4/data/nfl_games.csv')
```

I added the difference in score to see if that could introduce a relationship.

```
[3]: data['score_diff'] = (data['score1'] - data['score2'])
```

```
[4]: data['team1_win'] = np.where(data['score_diff'] > 0, 1, 0)
data['team2_win'] = np.where(data['score_diff'] < 0, 1, 0)
```

```
[5]: #data['date'] = pd.to_datetime(data['date'], format='%Y-%m-%d')
data['elo_diff'] = (data['elo1'] - data['elo2'])
```

```
[6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16274 entries, 0 to 16273
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date         16274 non-null  object
1   season       16274 non-null  int64
2   neutral      16274 non-null  int64
```

```

3  playoff      16274 non-null  int64
4  team1        16274 non-null  object
5  team2        16274 non-null  object
6  elo1         16274 non-null  float64
7  elo2         16274 non-null  float64
8  elo_prob1    16274 non-null  float64
9  score1       16274 non-null  int64
10 score2       16274 non-null  int64
11 result1     16274 non-null  float64
12 score_diff  16274 non-null  int64
13 team1_win   16274 non-null  int64
14 team2_win   16274 non-null  int64
15 elo_diff    16274 non-null  float64
dtypes: float64(5), int64(8), object(3)
memory usage: 2.0+ MB

```

```
[7]: data = data.dropna()
```

```
[8]: data.describe()
```

```

[8]:
      count      season      neutral      playoff      elo1      elo2  \
count  16274.000000  16274.000000  16274.000000  16274.000000  16274.000000  16274.000000
mean    1982.437569     0.005223     0.034779    1502.458394    1498.918375
std       25.448049     0.072084     0.183226     105.015371     104.541271
min     1920.000000     0.000000     0.000000    1119.595000    1156.551000
25%     1967.000000     0.000000     0.000000    1429.242750    1425.864750
50%     1987.000000     0.000000     0.000000    1504.015000    1500.185000
75%     2003.000000     0.000000     0.000000    1578.071500    1575.753000
max     2018.000000     1.000000     1.000000    1839.663000    1849.484000

      elo_prob1      score1      score2      result1      score_diff  \
count  16274.000000  16274.000000  16274.000000  16274.000000  16274.000000
mean     0.584829    21.544058    18.578161     0.580681     2.965897
std     0.175302    11.289422    10.794566     0.488551    15.469792
min     0.070953     0.000000     0.000000     0.000000    -73.000000
25%     0.461231    14.000000    10.000000     0.000000     -7.000000
50%     0.596354    21.000000    17.000000     1.000000     3.000000
75%     0.719930    28.000000    26.000000     1.000000    13.000000
max     0.970516    72.000000    73.000000     1.000000    66.000000

      team1_win      team2_win      elo_diff
count  16274.000000  16274.000000  16274.000000
mean     0.571034     0.409672     3.540020
std     0.494944     0.491788    141.791884
min     0.000000     0.000000   -511.826000
25%     0.000000     0.000000   -91.493250
50%     1.000000     0.000000     3.203000

```

75%	1.000000	1.000000	99.353500
max	1.000000	1.000000	541.969000

There are 16,274 entries. The first season starting in 1920. I agree that some teams that were good in 1920 might still be good today, but I think I will only take into account the last 15 years. So I will drop any game before the 2005 season.

```
[9]: data = data[data['season'] >= 2005]
```

```
[10]: data.head(5)
```

```
[10]:
```

	date	season	neutral	playoff	team1	team2	elo1	elo2	\
12536	2005-09-08	2005	0	0	NE	OAK	1712.670	1403.871	
12537	2005-09-11	2005	0	0	CAR	NO	1530.111	1506.978	
12538	2005-09-11	2005	0	0	KC	NYJ	1533.474	1547.067	
12539	2005-09-11	2005	0	0	JAX	SEA	1485.276	1480.542	
12540	2005-09-11	2005	0	0	WSH	CHI	1459.026	1410.741	

	elo_prob1	score1	score2	result1	score_diff	team1_win	team2_win	\
12536	0.895833	30	20	1.0	10	1	0	
12537	0.624181	20	23	0.0	-3	0	1	
12538	0.573445	27	7	1.0	20	1	0	
12539	0.599029	26	14	1.0	12	1	0	
12540	0.657488	9	7	1.0	2	1	0	

	elo_diff
12536	308.799
12537	23.133
12538	-13.593
12539	4.734
12540	48.285

```
[11]: data.team1.unique()
```

```
[11]: array(['NE', 'CAR', 'KC', 'JAX', 'WSH', 'SF', 'CLE', 'BUF', 'LAC', 'MIN',
        'MIA', 'DET', 'PIT', 'BAL', 'NYG', 'ATL', 'CIN', 'TB', 'TEN',
        'OAK', 'CHI', 'NYJ', 'DEN', 'HOU', 'GB', 'IND', 'ARI', 'SEA',
        'PHI', 'DAL', 'LAR', 'NO'], dtype=object)
```

By dropping games before 2005, I now have a better idea of what teams are currently active.

2 Faceting

I decided to try faceting to better understand the data.

Tips for faceting I found at <https://www.kaggle.com/residentmario/faceting-with-seaborn>


I was interested in comparing scores and the teams that played.

```
[12]: import seaborn as sns
```

```
[13]: # data2010 = data[data['season'].isin(['2010'])]
```

```
[14]: sns.pairplot(data,  
    x_vars = ['elo1', 'elo2', 'elo_prob1', 'team1_win', 'team2_win'],  
    y_vars = ['score_diff', 'score1', 'score2'])  
    #hue = "team1")
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x11ed350d0>
```



It looks like there may be a linear correlation with score1 and elo1 and score2 and elo2. Although that seems intuitive.

I didn't see any great correlations to use. So, I moved onto a heat map. The closer to 1 or -1 the value is on the map, the stronger the correlation. There isn't a great correlation between anything except maybe elo_prob1 and result_1. That seems like the highest non-trivial value.

```
[15]: data_corr = data.corr()  
data_corr
```

```
[15]:
```

	season	neutral	playoff	elo1	elo2	\
season	1.000000e+00	0.028296	-5.931427e-19	0.007367	0.005822	
neutral	2.829635e-02	1.000000	1.499114e-01	0.041492	0.036028	
playoff	-5.931427e-19	0.149911	1.000000e+00	0.264268	0.241292	
elo1	7.366663e-03	0.041492	2.642684e-01	1.000000	0.074469	


elo2	5.821616e-03	0.036028	2.412920e-01	0.074469	1.000000
elo_prob1	2.189256e-03	-0.049735	1.789909e-02	0.675980	-0.673994
score1	6.174932e-02	0.001141	3.029494e-02	0.250805	-0.124768
score2	6.417428e-02	0.002981	7.697921e-03	-0.153109	0.219881
result1	9.271437e-03	0.003245	1.784226e-02	0.237555	-0.209652
score_diff	-1.188288e-04	-0.001217	1.603222e-02	0.279407	-0.236212
team1_win	7.715245e-03	0.000999	1.826544e-02	0.237385	-0.209955
team2_win	-1.080838e-02	-0.005485	-1.737979e-02	-0.237205	0.208890
elo_diff	1.163446e-03	0.004179	1.795476e-02	0.682533	-0.677998

	elo_prob1	score1	score2	result1	score_diff	team1_win \
season	0.002189	0.061749	0.064174	0.009271	-0.000119	0.007715
neutral	-0.049735	0.001141	0.002981	0.003245	-0.001217	0.000999
playoff	0.017899	0.030295	0.007698	0.017842	0.016032	0.018265
elo1	0.675980	0.250805	-0.153109	0.237555	0.279407	0.237385
elo2	-0.673994	-0.124768	0.219881	-0.209652	-0.236212	-0.209955
elo_prob1	1.000000	0.275833	-0.271637	0.327205	0.377132	0.327480
score1	0.275833	1.000000	-0.053653	0.557326	0.737383	0.556660
score2	-0.271637	-0.053653	1.000000	-0.571707	-0.714065	-0.571322
result1	0.327205	0.557326	-0.571707	1.000000	0.777466	0.998908
score_diff	0.377132	0.737383	-0.714065	0.777466	1.000000	0.776739
team1_win	0.327480	0.556660	-0.571322	0.998908	0.776739	1.000000
team2_win	-0.326215	-0.556774	0.570843	-0.998907	-0.776495	-0.995632
elo_diff	0.992235	0.276312	-0.274007	0.328756	0.379071	0.328853

	team2_win	elo_diff
season	-0.010808	0.001163
neutral	-0.005485	0.004179
playoff	-0.017380	0.017955
elo1	-0.237205	0.682533
elo2	0.208890	-0.677998
elo_prob1	-0.326215	0.992235
score1	-0.556774	0.276312
score2	0.570843	-0.274007
result1	-0.998907	0.328756
score_diff	-0.776495	0.379071
team1_win	-0.995632	0.328853
team2_win	1.000000	-0.327940
elo_diff	-0.327940	1.000000

```
[16]: sns.heatmap(data_corr, vmin=-1, vmax=1, cmap="viridis", annot=True, linewidth=0.
      ↪ 1)
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x11af6b940>
```



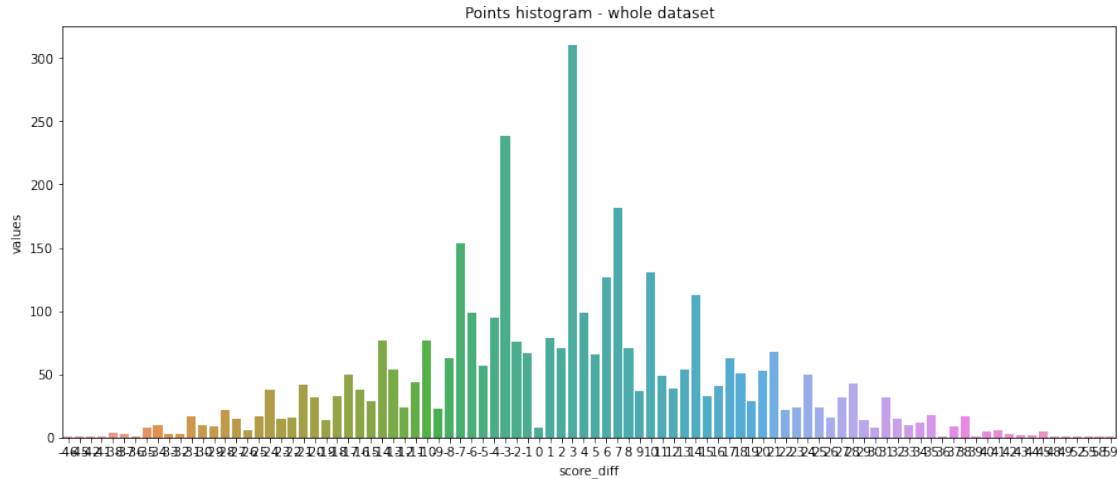
```
[17]: import matplotlib.pyplot as plt

def pastel_plot(data, x, y):
    plt.figure(figsize = (15,6))
    plt.title('Points histogram - whole dataset')
    sns.set_color_codes("pastel")
    sns.barplot(x = x, y=y, data=df)
    locs, labels = plt.xticks()
    plt.show()
```


I thought maybe a histogram would help to look at the distribution of the values for the difference in score. This revealed there were no crazy outliers.

```
[18]: temp = data['score_diff'].value_counts()
df = pd.DataFrame({'score_diff': temp.index,
                  'values': temp.values
                  })

pastel_plot(df, 'score_diff', 'values')
```




```
[19]: plt.figure(figsize=(20,5))
plt.title("Distribution of score")
ax = sns.distplot(data["score_diff"])
```



```
[20]: data=data.sort_values('score_diff', ascending=False)
```

The following bar graphs revealed the best teams in the team1 category and the worst teams in the team2 category.


```
[21]: plt.figure(figsize = (14,6))
plt.title('scores for different teams')
sns.barplot(x = 'team1', y="score_diff", data=data.head(20))
locs, labels = plt.xticks()
plt.show()
```



NE has the greatest score difference. This means they win by the most amount of points. Still, I needed to figure out some way to interact the two teams.

Also, below TEN loses by the most amount of points.

```
[22]: plt.figure(figsize = (14,6))
plt.title('scores for different teams')
sns.barplot(x = 'team2', y="score_diff", data=data.head(20))
locs, labels = plt.xticks()
plt.show()
```



3 Feature Engineering

Before I could continue, I needed some way to get an interaction between the two teams. I started by making a one hot encoding of team1 and team2. I then combined the encodings. To maintain a sense of who won and who lost, I made the team2 values negative, ie -1, and the team1 values positive. Then, if the score difference was negative, team2 won and if the score difference was positive, team1 won.

```
[23]: from random import shuffle
      data = data.sample(frac = 1)
```

```
[24]: team1_oh = pd.get_dummies(data['team1'], dtype=np.int64)
      team2_oh = pd.get_dummies(data['team2'], dtype=np.int64)
```

```
[25]: team1_oh.head(5)
```

```
[25]:
```

	ARI	ATL	BAL	BUF	CAR	CHI	CIN	CLE	DAL	DEN	...	NYG	NYJ	OAK	\
12954	0	0	1	0	0	0	0	0	0	0	...	0	0	0	
13206	0	0	1	0	0	0	0	0	0	0	...	0	0	0	
15459	0	0	0	0	0	0	0	0	0	1	...	0	0	0	
14659	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
16075	0	0	0	0	0	0	1	0	0	0	...	0	0	0	

	PHI	PIT	SEA	SF	TB	TEN	WSH
12954	0	0	0	0	0	0	0
13206	0	0	0	0	0	0	0
15459	0	0	0	0	0	0	0
14659	0	0	0	0	0	1	0
16075	0	0	0	0	0	0	0

[5 rows x 32 columns]

Now, I need to fit the teams together.

```
[26]: teams_oh = team1_oh.sub(team2_oh)
      teams_oh['score_diff'] = data['score_diff']
      teams_oh['elo_diff'] = data['elo_diff']
      teams_oh['date'] = data['date']
      teams_oh.head(5)
```

```
[26]:
```

	ARI	ATL	BAL	BUF	CAR	CHI	CIN	CLE	DAL	DEN	...	PHI	PIT	SEA	\
12954	0	-1	1	0	0	0	0	0	0	0	...	0	0	0	
13206	0	0	1	0	0	0	-1	0	0	0	...	0	0	0	
15459	0	0	0	0	0	0	0	0	0	1	...	0	0	0	
14659	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
16075	0	0	0	0	0	0	1	0	0	0	...	0	0	0	

	SF	TB	TEN	WSH	score_diff	elo_diff	date
12954							
13206							
15459							
14659							
16075							

12954	0	0	0	0	14	114.856000	2006-11-19
13206	0	0	0	0	-14	78.454000	2007-11-11
15459	0	0	0	0	7	228.311337	2016-01-03
14659	0	0	1	0	18	137.787000	2012-12-30
16075	0	0	0	0	10	26.205895	2018-10-07

[5 rows x 35 columns]

If the row has a 1, it was team1 and if the row has a -1, then it was team2.

If the score is positive, then team1 won, and if the score was negative then team2 won.

Encoding this way ensures I can know which team won and which team lost.

Later in the project, I realized I needed to predict the scores of each team, not their difference. So I made the data table `scores` to be able to predict the score of the teams.

```
[27]: scores = data[['score1', 'score2']]
      scores.head()
```

```
[27]:      score1  score2
      12954      24      10
      13206       7      21
      15459      27      20
      14659      38      20
      16075      27      17
```

4 Regression

Regression is useful to determine the relationship between independent variables (features) and dependent variables (target values).

In this problem, I tasked to predict the scores based on features. That is, the score is a function of the teams playing, and the elo predictions.

4.1 Linear Regression

I used <https://www.geeksforgeeks.org/linear-regression-python-implementation/> for help in implementing linear regression.

This was my first attempt at regression. I tried to fit the data to predict the score difference. The coefficients of the regression were the rating of each time. If the team was a good, then they had a high positive rating. If the team was loosing, then the had a high negative rating.

From the faceting before, TEN should have a low negative rating and NE should have a high positive rating. This is true below.

```
[28]: from sklearn import datasets, linear_model, metrics
      from sklearn.linear_model import Ridge, LinearRegression
      import matplotlib.pyplot as plt
```

```
[29]: data_train = teams_oh
data_train = data_train.dropna(axis='columns')
data_train.head(5)
```

```
[29]:
```

	ARI	ATL	BAL	BUF	CAR	CHI	CIN	CLE	DAL	DEN	...	PHI	PIT	SEA	\
12954	0	-1	1	0	0	0	0	0	0	0	...	0	0	0	
13206	0	0	1	0	0	0	-1	0	0	0	...	0	0	0	
15459	0	0	0	0	0	0	0	0	0	1	...	0	0	0	
14659	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
16075	0	0	0	0	0	0	1	0	0	0	...	0	0	0	

	SF	TB	TEN	WSH	score_diff	elo_diff	date
12954	0	0	0	0	14	114.856000	2006-11-19
13206	0	0	0	0	-14	78.454000	2007-11-11
15459	0	0	0	0	7	228.311337	2016-01-03
14659	0	0	1	0	18	137.787000	2012-12-30
16075	0	0	0	0	10	26.205895	2018-10-07

[5 rows x 35 columns]

```
[30]: # defining feature matrix(X) and response vector(y)
X = data_train.drop(['score_diff'], axis=1)
X = data_train.drop(['date'], axis=1)
#X = np.array(data.elo1).reshape(-1, 1)
y = data_train['score_diff']
```

```
[31]: # splitting X and y into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=1)
```

```
[32]: lr = Ridge(alpha=0.001)
lr.fit(X_train, y_train)
```

```
[32]: Ridge(alpha=0.001)
```

```
[33]: df_ratings = pd.DataFrame(data={'team': X.columns, 'rating': lr.coef_})
df_ratings
```

```
[33]:
```

	team	rating
0	ARI	-2.532983e-09
1	ATL	1.041983e-09
2	BAL	4.207294e-09
3	BUF	-2.992732e-09
4	CAR	4.624595e-10
5	CHI	2.845380e-09
6	CIN	-1.355059e-09

```

7      CLE -2.302405e-09
8      DAL  2.169938e-09
9      DEN -9.807869e-10
10     DET -1.777190e-09
11     GB  3.438865e-09
12     HOU -8.266760e-10
13     IND  3.359043e-10
14     JAX -2.593188e-09
15     KC  6.747566e-10
16     LAC  4.960689e-09
17     LAR -2.454958e-09
18     MIA -2.852748e-09
19     MIN  1.227153e-10
20     NE  6.434558e-09
21     NO  3.456554e-09
22     NYG  1.861281e-09
23     NYJ -2.200696e-09
24     OAK -5.696832e-09
25     PHI  2.540132e-09
26     PIT  3.378906e-09
27     SEA  2.902558e-09
28     SF -1.823715e-09
29     TB -2.249039e-09
30     TEN -2.815024e-09
31     WSH -1.712477e-09
32 score_diff 1.000000e+00
33 elo_diff  5.629231e-11

```

```
[34]: print('Score: \n', lr.score(X_test, y_test))
```

```
Score:
1.0
```

As found here: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html

The score is the R^2 coefficient for this fit. A model is “good” when the R^2 value is close to 1. A score of exactly 1 means the model is overfit. It’s not a good representation of the data.


The graph below showing the residual errors also proves this is not a good model.

```
[35]: plt.style.use('fivethirtyeight')
plt.scatter(lr.predict(X_train), lr.predict(X_train) - y_train, color = 'green', s = 10, label = 'Train data')
plt.scatter(lr.predict(X_test), lr.predict(X_test) - y_test, color = "blue", s = 10, label = 'Test data')
## plotting line for zero residual error
plt.hlines(y = 0, xmin = 0, xmax = 50, linewidth = 2)
## plotting legend
plt.legend(loc = 'upper right')
```

```

## plot title
plt.title("Residual errors")
## function to show plot
plt.show()

```



It is overfit.

4.2 Random Forest

I tried again with Random Forest Regression. This time, instead of predicting the score difference, I predicted the scores of each team using the `scores` data set.

Here is where I got tips to do it <https://towardsdatascience.com/random-forest-and-its-implementation-71824ced454f>

```

[36]: #X = teams_oh.drop(['score_diff'], axis=1)
      X = teams_oh.drop(['date'], axis=1)
      y = scores

```

```

[37]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=1)

```

```

[38]: X_train

```

```
[38]:
```

	ARI	ATL	BAL	BUF	CAR	CHI	CIN	CLE	DAL	DEN	...	OAK	PHI	PIT	\
15729	0	0	0	0	0	0	0	0	0	0	...	-1	0	0	
15712	0	0	0	0	0	0	0	0	1	0	...	0	0	0	
14405	0	0	0	0	0	0	0	0	-1	0	...	0	0	0	
13614	0	0	0	0	1	0	0	0	0	0	...	0	-1	0	
15517	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
...
15683	0	0	0	0	0	0	1	0	0	0	...	0	0	-1	
13768	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
13201	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
15510	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
13524	0	0	0	0	0	0	0	0	0	0	...	0	0	0	

	SEA	SF	TB	TEN	WSH	score_diff	elo_diff
15729	0	0	0	0	0	13	-90.945651
15712	0	0	0	0	0	21	84.793225
14405	0	0	0	0	0	-7	125.441000
13614	0	0	0	0	0	-28	-74.013000
15517	0	0	0	0	0	21	66.198425
...
15683	0	0	0	0	0	-4	-86.545743
13768	0	0	0	0	0	29	247.249000
13201	0	0	0	0	0	2	-153.999000
15510	0	0	0	0	-1	-2	40.148062
13524	0	0	1	0	0	3	48.438000

[2990 rows x 34 columns]

Number of estimators is equal to the number of trees.

```
[39]: from sklearn.ensemble import RandomForestRegressor
regressor = RandomForestRegressor(n_estimators = 10, random_state = 0)
regressor.fit(X, y)
```

```
[39]: RandomForestRegressor(n_estimators=10, random_state=0)
```

```
[40]: pred = regressor.predict(X_test)
```

```
[41]: print(metrics.r2_score(y_test, pred))
```

0.9058527703721649

This model has an R^2 value closer to 1. The value is 0.82. I am getting closer!

```
[42]: from sklearn.ensemble import RandomForestRegressor
regressor1 = RandomForestRegressor(n_estimators = 100, random_state = 0)
regressor1.fit(X, y)
pred1 = regressor1.predict(X_test)
```

```
print(metrics.r2_score(y_test, pred1))
```

0.928080229287205

It seems that more trees provide a better value for R^2 .

```
[43]: from sklearn.ensemble import RandomForestRegressor
regressor = RandomForestRegressor(n_estimators = 500, random_state = 0)
regressor.fit(X, y)
pred = regressor.predict(X_test)
print(metrics.r2_score(y_test, pred))
```

0.9310127291918824

I needed to do cross-validation to test if the model was a good fit.

I learned how to do cross validation here: https://jamesrledoux.com/code/k_fold_cross_validation

```
[44]: from sklearn.model_selection import cross_validate
cv = cross_validate(regressor1, X_test, y_test, cv=5)
print(cv['test_score'])
print(cv['test_score'].mean())
```

[0.52270444 0.39376462 0.41070127 0.32560158 0.42161389]
0.414877160534797

So, I see that this is not a good model at all. Ideally, you would want the values close to 0.99 to get an accurate fit of the data. Nevertheless, I proceeded with this choice to predict scores.

5 Predicting Scores

Now that I have the test data, I used the model to predict the scores between some teams.

I first turned the data into a dataframe.

```
[45]: pred_scores = pd.DataFrame({'predscore_1': pred1[:, 0], 'predscore_2': pred1[:, 1]})
pred_scores
```

```
[45]:
```

	predscore_1	predscore_2
0	18.19	29.28
1	22.61	16.59
2	27.27	14.05
3	23.84	29.78
4	40.96	9.16
..
743	28.53	10.32
744	24.86	10.90
745	30.57	27.59
746	21.56	29.39

747 53.91 5.98

[748 rows x 2 columns]

```
[46]: X_test['team1']=''  
      X_test['team2']=''
```

<ipython-input-46-aace447a7dd7>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
      X_test['team1']=''
```

<ipython-input-46-aace447a7dd7>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
      X_test['team2']=''
```

I readdd which team is which.

```
[47]: for col_name in X_test.columns:  
      #X_test.loc[X_test[col_name]==1, 'team1']= X_test['team1']+' '+col_name  
      X_test.loc[X_test[col_name]==1, 'team1']= col_name  
      X_test.loc[X_test[col_name]==-1, 'team2']= col_name  
  
      X_test = X_test.reset_index()  
      X_test = X_test.drop(columns='index')  
      X_test
```

/Users/annarosefritz/opt/anaconda3/lib/python3.8/site-packages/pandas/core/indexing.py:966: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
      self.obj[item] = s
```

```
[47]:
```

	ARI	ATL	BAL	BUF	CAR	CHI	CIN	CLE	DAL	DEN	...	PIT	SEA	SF	TB	\
0	1	0	0	0	0	0	0	0	-1	0	...	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	
2	0	0	0	0	-1	0	0	0	0	0	...	0	0	0	0	
3	0	0	0	0	0	0	0	0	0	0	...	1	0	0	0	
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	


```

..    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ..
743    0    0    1    0    0    0    0    -1    0    0    0    ...    0    0    0    0
744    0    0    0    0    0    0    0    0    0    0    0    ...    0    1    0    0
745    0    0    0    0    0    0    0    0    0    0    0    ...    0    -1    0    0
746    0    0    0    0    0    0    0    -1    0    0    0    ...    0    0    0    0
747    0    0    0    0    0    0    0    0    0    0    0    ...    0    0    0    0

```

```

      TEN  WSH  score_diff  elo_diff  team1  team2
0        0    0          -11 -31.646891   ARI   DAL
1        1    0           6  68.641000   TEN   JAX
2        0    0          13 -65.287000  NYG   CAR
3        0    0          -6 -49.732000  PIT    GB
4        0    0          32  63.115000  MIN   DET
..    ...    ...    ...    ...    ...    ...
743    0    0           18   9.831000  BAL   CLE
744    0    0           14 -76.314334  SEA   PHI
745    0    0           3 -200.607000  LAR   SEA
746    1    0          -8 -108.414000  TEN   CIN
747   -1    0          48  237.609000   GB   TEN

```

[748 rows x 36 columns]

```

[48]: df_pred_scores = X_test[['team1', 'team2']]
df_pred_scores = pd.merge(df_pred_scores, pred_scores, on=df_pred_scores.index)
df_pred_scores['home_away'] = list(zip(df_pred_scores['team1'],
    ↪df_pred_scores['team2']))
df_pred_scores

```

```

[48]:   key_0  team1  team2  predscore_1  predscore_2  home_away
0        0   ARI   DAL         18.19         29.28  (ARI, DAL)
1        1   TEN   JAX         22.61         16.59  (TEN, JAX)
2        2  NYG   CAR         27.27         14.05  (NYG, CAR)
3        3  PIT   GB          23.84         29.78  (PIT, GB)
4        4  MIN   DET         40.96          9.16  (MIN, DET)
..    ...    ...    ...    ...    ...    ...
743    743  BAL   CLE         28.53         10.32  (BAL, CLE)
744    744  SEA   PHI         24.86         10.90  (SEA, PHI)
745    745  LAR   SEA         30.57         27.59  (LAR, SEA)
746    746  TEN   CIN         21.56         29.39  (TEN, CIN)
747    747   GB   TEN         53.91          5.98  (GB, TEN)

```

[748 rows x 6 columns]

```

[49]: def predict (team1_in, team2_in):

      df = df_pred_scores[(df_pred_scores['team1'] == team1_in) &
    ↪(df_pred_scores['team2'] == team2_in)]

```

```
print(team1_in, "score is predicted to be:", df[['predscore_1']])  
print(team2_in, "score is predicted to be:", df[['predscore_2']])
```

```
[50]: predict('PIT', 'KC')
```

```
PIT score is predicted to be:    predscore_1  
351          40.83  
KC score is predicted to be:    predscore_2  
351          11.84
```

Here it gives you the predicted score of the game.